

Distributed & network programming (2021 Fall)

Lab 01: Sequential (serial) vs Parallel processing



Aug 16, 2021

Prof. Shinnazar Seytnazarov

Faculty of Computer Science & Engineering

Outline

- ❑ **Introduction to process and multiprocessing in Python**
- ❑ **Exercise #1: Sequential processing (20pts)**
- ❑ **Exercise #2: Parallel processing (40pts)**
 - Using multiprocessing library in Python
- ❑ **Exercise #3: Comparison of Sequential and Parallel (40pts)**
 - Drawing plot for different number of processes
 - Explain the results

Traditional processes


❑ A process is

- an instance of a computer program that is being executed.

❑ A process has three basic components:

- An executable program
- The associated resources/data needed by the program (variables, workspace, buffers, etc.)
- The execution context (state of process)

The process scheduling in single core CPU

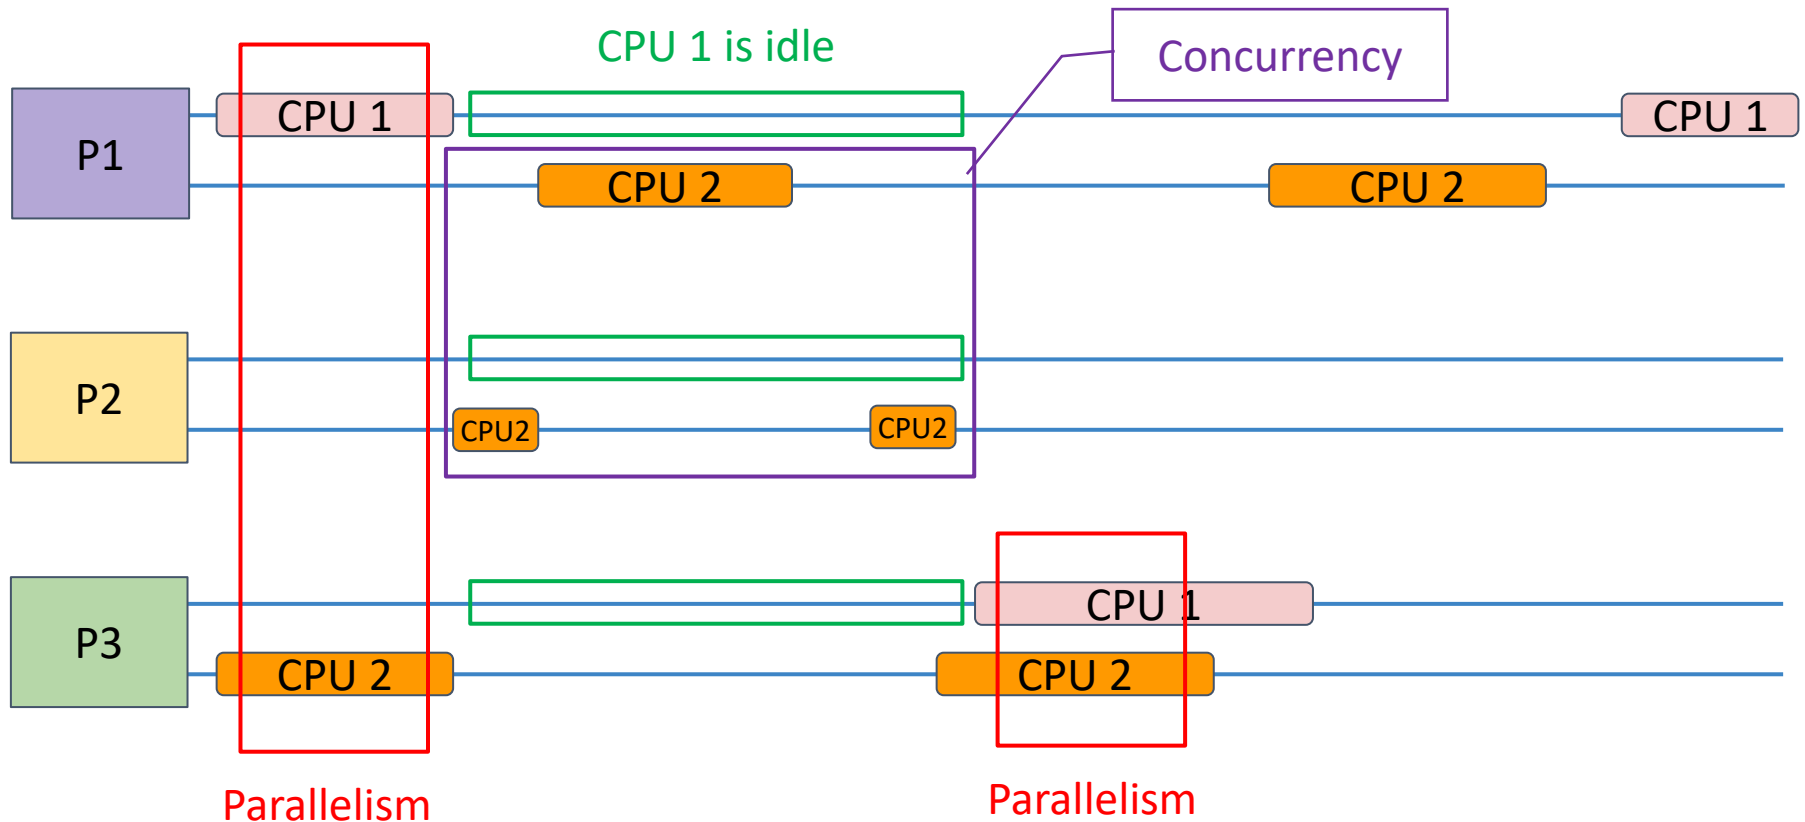


This is what we call:

Concurrency

Several processes are “running” but they share the same CPU.

A two-core CPU



- ❑ **Concurrency** - Multiple processes share the CPU
- ❑ **Parallelism** - Multiple processes are executed simultaneously on different CPUs on parallel

Parallel processing in Python

❑ Supported by *multiprocessing* module

❑ `Process(group=None, target=None, name=None, args=(), kwargs={}, daemon=None)`

- This constructor should always be called with **keyword arguments**.
- **group** should be None; reserved for future extension.
- **target** is the callable object to be invoked by the `run()` method.
- **name** is the process name. By default, a unique name is constructed of the form “process-N” where N is a small decimal number.
- **args** is the argument tuple for the target invocation.
- **kwargs** is a dictionary of keyword arguments for the target invocation.
- If not None, **daemon** explicitly sets whether the thread is daemonic. If None (the default), the daemonic property is inherited from the current process.

Process class methods

❑ **start()**

- Start the process' activity

❑ **is_alive()**

- Return whether the process is alive.

❑ **join(timeout=None)**

- Wait until the process terminates. This blocks the calling process until the process whose join() method is called terminates

❑ **run()**

- Method representing the thread's activity.
- You may override this method in a subclass. The standard **run()** method invokes the **target** argument.

Creating the processes

❑ Creating the processes

- Need to call the constructor of Process class with **keyword arguments**

- Creating a single thread
- Need to call **start()** method to run the thread

```
1 from multiprocessing import Process
2 from time import sleep, time
3
4 def greet(name):
5     print(f"{time()}: Hi {name}")
6     sleep(5)
7     print(f"{time()}: Bye {name}")
8
9 p = Process(target=greet, args=("Ivan",))
10 p.start() # starts the process's activity
11 p.join() # blocks the interpreter until t
12 p.close() # closes the Process object, re
```

```
1628930659.8606088: Hi Ivan
1628930664.9100642: Bye Ivan
```

- Creating multiple threads

```
1 processes = [Process(target=greet, args=(name,))
2               for name in ["Ivan", "Alexey", "Vladimir"]]
```

```
1 [p.start() for p in processes]
2 [p.join() for p in processes]
3 [p.close() for p in processes]
```

```
1628930793.4104717: Hi Ivan
1628930793.4286308: Hi Alexey
1628930793.449318: Hi Vladimir
1628930798.4274025: Bye Ivan
1628930798.452116: Bye Alexey
1628930798.4601228: Bye Vladimir
```


Effect of join() call

❑ join() blocks the main program

```
1 p = Process(target=greet, args=("Ivan",))
2 p.start()
3 sleep(1)
4 print("main program isn't blocked")
```

```
1628933949.9598112: Hi Ivan
main program isn't blocked
1628933954.9769063: Bye Ivan
```

```
1 p.close()
```


```
1 p = Process(target=greet, args=("Ivan",))
2 p.start()
3 p.join()
4 print("main program is blocked")
5 p.close()
```

```
1628933998.7069: Hi Ivan
1628934003.7225947: Bye Ivan
main program is blocked
```

Producer-Consumer model

□ Producer-Consumer model

- some processes produce tasks to do and put them in a shared data structure
- while some other processes "consume" the tasks from shared data structure and do the work.



A shared data structure: *Queue* class

- ❑ We can use the *Queue* class from multiprocessing module, as our shared data structure:
 - a thread safe FIFO (first in first out) queue.
 - The **queue** module implements multi-producer, multi-consumer queues.
 - It is especially useful in threaded programming when information must be exchanged safely between multiple processes.
 - The **Queue** class in this module implements all the required locking semantics.

Queue class

❑ Queue(maxsize=0)

- Constructor for a FIFO queue.
- **maxsize** is an integer that sets the upperbound limit on the number of items that can be placed in the queue.
- Insertion will block once this size has been reached, until queue items are consumed.
- If **maxsize** is less than or equal to zero, the queue size is infinite.

```
1 from multiprocessing import Queue
1 q = Queue()
1 q.qsize()
0
1 q.empty()
True
```

Some methods of Queue class

❑ `put(item, block=True, timeout=None):`

- Put item into the queue.
- If optional **block** is True and **timeout** is None (the default), blocks until a free slot is available.
- If **timeout** is a positive number, it blocks at most **timeout** seconds and raises the **Full** exception if no free slot was available within that time.
- If **block** is false, put an item on the queue if a free slot is immediately available, else raise the **Full** exception (timeout is ignored in that case).

1	<code>q = Queue(maxsize=3)</code>
1	<code>q.put("A")</code>
2	<code>q.qsize()</code>
1	
1	<code>q.empty()</code>
	False
1	<code>q.full()</code>
	False
1	<code>q.put("B")</code>
2	<code>q.put("C")</code>
1	<code>q.qsize()</code>
	3
1	<code>q.full()</code>
	True
1	<i># blocks the interpreter</i>
2	<i># when queue is full</i>
3	<i># blocks until item is put</i>
4	<code>q.put("D")</code>

Switching to non-blocking mode

❑ `put()` blocks the main program if the queue is full

- To overcome this, disable blocking mode
- And catch the `Full` exception to prevent the program to stop

```
1 q.full()
True

1 q.put("D", block=False)
-----
Full
<ipython-input-10-7641d34a6de2>
----> 1 q.put("D", block=False)

/usr/lib/python3.8/multiprocessing/queues.py:82: raise ValueError
82         raise ValueError('Queue is full')
83     if not self._sentinel:
--> 84         raise Full
85
86     with self._not_full:

Full:

1 from queue import Full

1 try:
2     q.put("D", block=False)
3 except Full:
4     print("Queue is full")

Queue is full
```

Some methods of Queue class

❑ **get(block=True, timeout=None)**

- Remove and return an item from the queue.
- If optional block is true and timeout is None (the default), block if necessary until an item is available.
- If timeout is a positive number, it blocks at most timeout seconds and raises the Empty exception if no item was available within that time.
- Otherwise (block is false), return an item if one is immediately available, else raise the Empty exception (timeout is ignored in that case).

```
1 q.get()
'A'

1 q.get()
'B'

1 q.qsize()
1

1 q.get()
'C'

1 q.empty()
True

1 # blocks the main program
2 q.get()

1 from queue import Empty
2
3 try:
4     q.get(block=False)
5 except Empty:
6     print("Queue is empty")

Queue is empty
```

Practice session

❑ Download the Lab01_assignment.ipynb file and solve the problems there

- You can also use .pdf file

❑ Submit your solution in jupyter-notebook (.ipynb) format

- Jupyter-notebook on Windows doesn't fully support Python's multiprocessing module
 - but it does fully support multiprocessing on Linux
 - Not sure about MacOS
- If you don't really want to work on Linux for some reason
 - solve the problems on one of the Windows IDEs: ex) vscode fully support multiprocessing module on Windows
 - then copy the code into jupyter-notebook and submit.

❑ Deadline 2021.08.22 PM11:59

- For each day of late submission, 5% will be cut
 - Ex) You submitted the next day after deadline, i.e., a day later, suppose you scored 70%, then your final score will be 65%.

Resources

- ❑ <https://docs.python.org/3/library/multiprocessing.html>
- ❑ <https://docs.python.org/3/library/multiprocessing.html#the-process-class>
- ❑ <https://docs.python.org/3/library/multiprocessing.html#pipes-and-queues>

Any questions?